Alex Morgan
Email: alex.morgan@email.com | Phone: +1-415-555-0182
LinkedIn: https://linkedin.com/in/alexmorgan-ai | GitHub: https://github.com/alexmorgan-ml
Location: San Francisco, CA

Professional Summary
Senior Machine Learning Engineer with 6+ years of specialized experience in natural language processing (NLP), transformer architectures, and generative AI systems. Proven record of developing scalable semantic search engines, fine-tuning large language models (LLMs), and deploying low-latency RAG systems serving over 2 million monthly queries.

Technical Skills
? AI & Machine Learning: PyTorch, Hugging Face, Transformers, NLP, Large Language Models (LLM), RAG, Fine-tuning, Vector Embeddings, Scikit-Learn
? Databases & Vector Search: PostgreSQL, ChromaDB, Pinecone, Redis
? Backend & Cloud: Python, FastAPI, Docker, Kubernetes, AWS, CI/CD, Git, Linux
? Methodologies: MLOps, System Design, Unit Testing, Agile

Work Experience
Senior Machine Learning Engineer | CogniWave AI (2022 - Present)
? Architected and deployed an enterprise Retrieval-Augmented Generation (RAG) assistant leveraging PyTorch, Hugging Face Transformers, and Pinecone, reducing customer support escalations by 42%.
? Fine-tuned open-source 7B parameter models using LoRA and QLoRA for domain-specific contract analysis with 91% F1 score.
? Containerized multi-agent inference microservices with Docker and deployed onto Kubernetes on AWS with horizontal auto-scaling.
? Implemented high-throughput asynchronous REST APIs using FastAPI, maintaining sub-120ms P99 latency under heavy traffic.

Machine Learning Engineer | Apex Data Labs (2019 - 2022)
? Built text classification and named entity recognition (NER) pipelines using Python, spaCy, and Scikit-Learn.
? Designed automated MLOps pipelines using MLflow, Docker, and GitHub Actions for continuous model evaluation and deployment.
? Integrated PostgreSQL and Redis caching layers to speed up real-time recommendation querying by 3x.

Key Projects
? OpenRAG Studio: Open-source toolkit for benchmarking dense vector search vs hybrid BM25 retrieval across multi-million document corpora. (Python, PyTorch, ChromaDB, FastAPI)
? BioMed-Summarizer: Clinical text abstractive summarization service fine-tuned on PubMed datasets using Hugging Face Transformers.

Education
Master of Science in Computer Science (Artificial Intelligence Specialization)
Stanford University | 2017 - 2019
Bachelor of Science in Computer Engineering
University of California, Berkeley | 2013 - 2017